Automated analysis of algorithms implemented on top of TBD

Walter Tichy, Umut Acar, Emanuel Jöbstl
Karlsruhe Institute of Technology, Faculty of Computer Science
Carnegie Mellon University, Computer Science Department

***Abstract*—The principle of incremental programming has been well known for many years. While the automatic transformation of sequential programs into efficient incremental programs is not solved in general, a number of platforms for incremental computation have been created. In this document, we inspect the incremental computation platform TBD. First, we summarize known principles like directed dependency graphs, execution traces, intrinsic trace distance and trace stability, then we create an program model which enables us to apply those principles to TBD programs. We also present an algorithm to calculate the intrinsic trace distance in practice. Finally, we show how we can make automatic optimizations to programs by analyzing the dependency graph.**

I. Introduction

This section is going to describe the purpose of incremental computation and also explain why the concept of incremental computation is useful for speeding up computations. Furthermore, used and referenced terminology should be outlined.

A. Approaches to incremental computation

This subsection briefly describes various approaches to incremental computation, and outline their strengths and weaknesses.

These approaches include:

- Providing a platform or framework for incremental programming, utilizing
  - function caching. [1] [2]
  - formal manipulation of the program. [3]
  - differential data flow. [4]
  - a combination of multiple approaches, like memorization and execution traces. [5] [6] [7] [8]
- Providing a high-level abstraction, like an incremental database. [9]
- Deriving an incremental program from a non-incremental, non-functional program. [10]
- Deriving an incremental program from a non-incremental, functional program. [11]

II. TBD

The TBD platform, a framework for incremental computation currently being developed at CMU. TBD follows the approach of memorization combined with directed dependency graphs (DDGs), as throughly described in [13]. Also, parallel computing is supported. The framework allows a programmer to write software using TBDs programming interface, while TBD automatically takes care about invoking the correct functions for change propagation, in case of an input data update. [**?**] The framework is being developed in the Scala language, which enables us to exploit the reflection capabilities of Scala for analysis [12] [**?**]. The source code of TBD is available at https://github.com/twmarshall/tbd.

```
def mod[T](
  initializer: Dest[T] => Changeable[T]
): Mod[T]
```

Fig. 1. Signature of the mod method

```
def write[T](
    dest: Dest[T],
    value: T
): Changeable[T]
```

Fig. 2. Signature of the $write$ method

A. Programming interface

TBD needs to keep track of function calls and reads and writes of variables in the program. To accomplish this, TBD wraps all values relevant for change propagation into so called *modifiables* or short *Mod*. TBD automatically wraps all input data into Mods.

1) mod: To create Mods, for example as result of the program execution, TBD provides a method mod. The declaration of mod can be seen in listing 1. The mod method calls a function parameter $initializer$ with a *destination* or *Dest* as argument. The value written to the Dest by the function parameter is then stored in the Mod, which is returned by the mod method. Forcing the return type of *Changeable* simply enforces that a write happens inside $initializer$.

2) write: To write to a Dest, TBD provides a $write$ method. The signature of $write$ can be found in listing 2. The $write$ method simply takes a Dest and a value, and writes the value to the given Dest. The write method returns a Changeable.

3) read: The values from within modifiables have to be read explicitely. For this purpose, TBD provides a $read$ method, which accepts a Mod as parameter and then calls a function parameter $reader$ with the value of the Mod as first argument. The signature can be seen in in listing 3. For $read$, the function parameter $reader$ also has to return a Changeable.

```
def read[T, U <: Changeable[_]](
    mod: Mod[T],
    reader: T => U
): U
```

Fig. 3. Signature of the *read* method

```
def add(
    tbd: TBD,
    mod1: Mod[Int],
    mod2: Mod[Int]
): Mod[Int]) = {
    tbd.mod((dest: Dest[Int]) => {
        tbd.read(mod1)(v1 => {
            tbd.read(mod2)(v2 => {
                tbd.write(dest, v1 + v2)
            })
        })
    })
}
```

Fig. 4. A basic example, utilizing *read*, *write*, and *mod*

Reads without an enclosed write are not useful, since the *read* method my not modify values outside of it's scope.

Listing 4 shows a very simple example, which adds two Mods of type integer. First, *mod* is called to create a Dest for the result, then the values of $mod1$ and $mod2$ are read. The values of $mod1$ and $mod2$ are then added and written to *dest*. The nested pattern of the functions, which is easily noticeable, is typical for TBD.

Since all programs consist of *read*, *write* and *mod* functions, and all Modifiables have to be explicitely written, TBD is able to construct a DDG from monitoring the calls to the corresponding functions.

4) memo: As we already mentioned, TBD not only utilized DDGs, but also memorization. To accomplish memorization, tbd provides a method to create so-called *Lifts*, which in turn provide a method for memorization, *memo*. The *memo* method accepts a list of parameters, which are used to match this *memo* call and a function parameter $func$. A Lift can be described as memorization context. Calling *memo* with the same parameters as any previous call on the same Lift will yield the same result, without evaluation $func$. If there is no match, $func$ will be called and the result will be stored for future memorization. In general, it is important to not share Lift objects between unrelated function calls, but to preserve the same Lift for all calls to the same function. The signature of *memo* can be seen in listing 5.

A typical use case for memorization is list processing. A typical example is shown in 6. First, we define a class for list nodes and the properties *value* of type integer and *next*. Note, that the class is immuteable. Next, we define a function, $incrementalList$, which initializes a lift and calls a recursive function, $incrementRecursive$ with the head of the list and the created lift. The latter function maps each list node to a list node with *value* increased by one. This is done by first creating a Dest *dest* for the new ListNode. Then, the current node is read from it's modifiable. If the current node is null, the end of the list is reached and *null* can be written to *dest*. If the current node is not *null*, the value is read, increased, and written again to create the Mod $newValue$, similar to the example in listing 4.

```
def memo(
    args: List[_],
    func: () => T
): T
```

Fig. 5. Signature of the *memo* method

```
def par[T, U](
    one: TBD => T,
    two: TBD => U
): Tuple2[T, U]
```

Fig. 7. Signature of the *par* method

Then, $incrementRecursive$ is called recursiveley with the next node as parameter. The call to $incrementRecursive$, however, is enclosed in a memo operation, with the next node as parameter. If now a change propagation happens, TBD is not going to recursiveley call all reads again, but will stop as soon as a memo match occours. This is typically the case as soon as the recursion reaches an unchanged list element.

In the end, a new list node is constructed from the results and returned.

5) par: The last crucial method offered by TBD is a method to execute code in parallel, *par*. The *par* method takes two function parameters *one* and *two*, where each function parameter is executed on a septerate worker thread with septerate TBD objects. The *par* method blocks util both workers are finished. The signature of *par* is shown in listing 7

B. Constraints and responsibilities

During change propagation, TBD re-evaluates all *read* calls that read modifiables which have changed, in the same order they were called during the initial run. Obviously, the functions invoked *read*, *mod*, *memo* and *par* may not write variables outside of their scope, or they will easily break change propagation. If, for example, a static variable is written from within a function called by read, and then used somewhere else in the program, the system has no way to propagate the change of this variable.

Furthermore, all functions called have to be deterministic. Calling the same function with the same parameters has to lead to the same return value or the same value written to a dest. Otherwise, memorization will not be useable in the program.

For each function parameter passed to the functions *read* or *mod*, the last operation executed in that function parameter has to be a *write*. This is enforced by requireing the return type of $Changeable$ for function parameters.

[Still, describe the order of function calls.]

```
class ListNode(_value: Mod[Int], _next: Mod[ListNode]) {
    val value = _value
    val next = _next
}

def incrementList(tbd: TBD, head: Mod[ListNode]): Mod[ListNode] = {
    val lift = tbd.makeLift()
    incrementRecursive(tbd, head, lift)
}

def incrementRecursive(tbd: TBD, current: Mod[ListNode], lift: Lift[ListNode])
    : Mod[ListNode] = {

    tbd.mod((dest: Dest[ListNode]) => {
        tbd.read(current)(current => {
            if(current == null) {
                tbd.write(dest, null)
            } else {
                val newValue = tbd.mod((destValue: Dest[Int]) => {
                    tbd.read(current.value)(value => {
                      tbd.write(destValue, value + 1)
                    })
                })

                val newNext = lift.memo(List(current.next), () => {
                    incrementRecursive(tbd, current.next, lift)
                })

                tbd.write(dest, new ListNode(newValue, newNext))
            }
        })
    })
}
```

Fig. 6. A basic example, utilizing $memo$

III. Theoretic fundamentals

The work of U. Acar et al[13] describes the theoretical concept of incremental computing using memorization and DDGs in detail. For our purpose, however, we have to adjust some of the fundamental definitions to match the TBD platform.

A. The Normal Form

B. Execution Traces

For completeness, we first This subsection describes the approaches of using Traces and Directed Dependency Graphs (DDGs). [13]

C. Memorization

This subsection describes how traces and memorization together are used to accomplish incremental computing. [13]

D. Stable algorithms

This subsection describes the concepts of stable algorithms, intrinsic trace distance and their relationship.

E. Intrinsic trace distance

[1]

Also, this section should emphasis that the intrinsic trace distance forms a lower bound for the time needed by change propagation during an update. [13]

F. Abstract machine model and normal form

IV. A program model for TBD

One of the base concepts described in the previous section is the concept of a trace. Atrace can be described as an ordered tree, whereas nodes represent function calls during the program execution. While we can retain the definition of a trace, we have to adjust the definition of nodes and node equality for our purpose. Also, we have to show that TBD programs fit the machine model described in[13].

A. Trace node equality and similarity

As described in section II, TBD provides$read$, mod, $write$, $memo$ and par methods to the developer. Instead of creating

[1] Intrinsic trace distance is a central concept for this work and can basically be described as an edit distance between two trees. The definition can be found in [13], chapter 7 or [14].

an execution trace out of all functions in the program, we restrict ourselfs to a trace consisting of only these functions. It should be noted, that, since we require each function to be side-effect free and determinisitc, we could theoretically omit $write$ nodes in the DDG, since they directly depend on their corresponding parent nodes. However, including these nodes can provide useful insights during debugging.

Definition 1 (Trace nodes) *Let each node in our execution trace represent a* $read$*,* mod*,* $write$ $memo$ *or* par *function. We annotate each node with a tuple of the following values:*

- *the node type* t*, which can have the values* $read$*,* mod*,* $write$*,* $memo$ *or* par
- *a node tag, a sequence of values which has a different structure depending on the node type*

Depending on the node type, we define the following node tags:

- *for* $read$ *nodes:*
  - *the value of the modifiable being read, denoted as* a
  - *the reader function being called, denoted as* fun
- *for* mod *nodes:*
  - *the initialzer function being called, denoted as* fun
- *for* $write$ *nodes:*
  - *the value being written, denoted as* a
  - *the destination being written to, denoted as* d
- *for* $memo$ *nodes:*
  - *the list of values to memo match against, denoted as* $(a_1, ..., a_n)$
  - *the function being called, denoted as* fun
- *for* par *nodes:*
  - *the first function being called, denoted as* fun_1
  - *the second function being called, denoted as* fun_2

[Todo: Include an example figure here.]

Given these definitions, we now re-define equality of nodes.

Definition 2 (Node equality) *Let a node* A *and* B *be equal, iff the node type of* A*,* t_a*, equals the node type of* B*,* t_b*, and the tag of* A *equals the tag of* B*.*

Obviously, we only compare the the tag if the node type already matches. Therefore, we can simply compare each element in the tag of A with it's counterpart in the tag of B.

The tag can consist of objects, value types, modifiables or functions. For comparing functions inside the tag we define equality as follows:

Definition 3 (Function execution equality for TBD traces) *A function execution* fun_a *and a function execution* fun_b *are equal, iff*

1) *they refer to the same symbol in the source code of the program*
2) *all arguments are equal*
3) *all free variables bound from an outer scope are equal.*

The requirement for side-effect free and deterministic functions leads to the conclusion, that all subcalls to other functions, including any writes, are going to be equal if the function is invoked with the same parameters. We have to take care of free variables in the function, however, since they might influence the behavior of the program. A prime example would be a $read$ nested within another $read$, whereas the inner $read$ accesses the value provided by the outer read, which can be seen in listing **??**. If the value of $mod1$ changes in the example the inner function performing the addition of $v1$ and $v2$ is not going to be equal anymore, therefore the $read$ node of the inner $read$ has changed, even the value of $mod2$ stays the same.

For comparing values or objects inside the tag, the function parameters or closed free variables we use $deepequality$. Modifiables, however should be compared by reference equality. The reason for doing so is to ensure correctness even with complex types, for example like arrays, nested lists or objects. For modifiables, the change propagation algorithm takes care of changed values, and automatically calls all subcalls which are affected. The case where the modifiable itself was re-created forms an exception, where we would have to re-execute all reads which would access this modifiable. This leads to the following formal definition:

Definition 4 (Object equality for TBD traces) *A pritive value* p *is equal to a primitive value* k *iff* p *and* k *have the same type and the same value.*

A modifiable x *is equal to a modifiable* y *iff* x *and* y *refer to the same object in memory.*

An object A *with ordered properties* $(a_1, ..., a_n)$ *is equal to an object* B *with ordered properties* $(b_1, ..., b_n)$ *iff* A *and* B *have the same type, and* a_i *equals* b_i $\forall i \in [1, n]$*. Properties can be other objects, modifiables or primitives.*

[Proof that equality/similarity definition is sufficient]

[Ask Umut wether trace vs. monotone trace, because if we have to keep ancestor relationship in mind, we HAVE to use LCS algo]

With these definition of trace node equality, we can keep the definition of $Cognates$ and $TraceDistance$ from [13].

B. TBD programs and the Normal Form

Before we can apply theorems regarding change propagation for TBD, we also have to show that TBD programs are in the normal form defined in [13].

V. An Intrinsic trace distance algorithm for TBD

This section is going to describe our algorithm for calculating the minimal trace distance of execution traces produced by TBD. Also, implementation details and challenges during the implementation (like defining equality for anonymous functions with free variables in practice) should be discussed.

A. Proof of correctness

We shall also proof the correctness of our algorithm. It is of importance that we take the changes of the program model outlined in the previous section into account.

VI. Automatic optimization of programs

This section is going to describe how analyzing the Directed Dependency Graph (DDG) can be used for automatic optimization. For accomplishing this task we can utilize the following features of the DDG:

- Caller/callee dependencies.
- Dependencies of modifiables[2].
- Dependencies of bound variables which are not modifiables.

Furthermore, using the intrinsic distance algorithm, we can recognize which nodes are deleted, inserted or retained [13], which can be used to optimize the program to accomplish faster change propagation.

The exact contents described in this chapter are still to be determined, based on our findings. Possible approaches include, but may not be limited to:

- Function call reordering.
- Insertion of explicit memorization calls.
- Detection of cascading updates, which could be omitted.

VII. Evaluation

This section is going to demonstrate the usefulness of the described techniques using real-world algorithms, like map, reduce and quicksort.

Basically, it is shown how it is possible to optimize a classic implementation (without memorization) of each algorithm, so that change propagation time lies within the same complexity class as the theoretical lower bound for updates for this algorithm.

VIII. Conclusion

This section will conclude and summarize with the findings of this work.

A. Future work

The final section briefly outlines problems encountered but not solved during the writing of this thesis, as well as encourages future research on interesting issues of incremental computation.

[2]Pointer-like variables which have to be explicitly read and written, and therefore support automatic change propagation

References

[1] A. Heydon, R. Levin, and Y. Yu, "Caching function calls using precise dependencies," *ACM SIGPLAN Notices*, vol. 35, no. 5, pp. 311–320, 2000.

[2] W. Pugh and T. Teitelbaum, "Incremental computation via function caching," in *Proceedings of the 16th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM, 1989, p. 315328. [Online]. Available: http://dl.acm.org/citation.cfm?id=75305

[3] R. F. Cohen and R. Tamassia, "Dynamic expression trees and their applications," in *SODA*, 1991, pp. 52–61.

[4] F. McSherry, R. Isaacs, M. Isard, and D. G. Murray, "Composable incremental and iterative data-parallel computation with naiad," Tech. Rep. MSR-TR-2012-105, October 2012. [Online]. Available: http://research.microsoft.com/apps/pubs/default.aspx?id=174076

[5] M. A. Hammer, U. A. Acar, and Y. Chen, "CEAL: a C-based language for self-adjusting computation," in *ACM Sigplan Notices*, vol. 44. ACM, 2009, p. 2537. [Online]. Available: http://dl.acm.org/citation.cfm?id=1542480

[6] Y. Chen, U. A. Acar, and K. Tangwongsan, "Functional Programming for Dynamic and Large Data with Self-Adjusting Computation," 2014. [Online]. Available: http://www.mpi-sws.org/~chenyan/papers/icfp14.pdf

[7] U. A. Acar, A. Ahmed, and M. Blume, "Imperative self-adjusting computation," in *ACM SIGPLAN Notices*, vol. 43. ACM, 2008, p. 309322. [Online]. Available: http://dl.acm.org/citation.cfm?id=1328476

[8] U. A. Acar, G. E. Blelloch, and R. Harper, "Adaptive functional programming," *ACM Transactions on Programming Languages and Systems (TOPLAS)*, vol. 28, no. 6, pp. 990–1034, 2006.

[9] D. Peng and F. Dabek, "Large-scale Incremental Processing Using Distributed Transactions and Notifications." in *OSDI*, vol. 10, 2010, p. 115. [Online]. Available: https://www.usenix.org/legacy/events/osdi10/tech/full_papers/Peng.pdf?origin=publication_detail

[10] Y. A. Liu and T. Teitelbaum, "Systematic derivation of incremental programs," *Science of Computer Programming*, vol. 24, no. 1, pp. 1–39, 1995.

[11] R. Ley-Wild, M. Fluet, and U. A. Acar, "Compiling self-adjusting programs with continuations," in *ACM Sigplan Notices*, vol. 43, no. 9. ACM, 2008, pp. 321–334.

[12] E. Burmako, "Scala macros: let our powers combine!: on how rich syntax and static types work with metaprogramming," in *Proceedings of the 4th Workshop on Scala*. ACM, 2013, p. 3.

[13] U. A. Acar, "Self-adjusting computation," 2005.

[14] U. A. Acar, G. E. Blelloch, R. Harper, J. L. Vittes, and S. L. M. Woo, "Dynamizing static algorithms, with applications to dynamic trees and history independence," in *Proceedings of the fifteenth annual ACM-SIAM symposium on Discrete algorithms*. Society for Industrial and Applied Mathematics, 2004, pp. 531–540.